

# Deep Learning for Natural Language Processing

## Conditional Random Fields



UNIVERSITY OF  
GOTHENBURG

---

**CHALMERS**

**WASP** | WALLENBERG AI  
AUTONOMOUS SYSTEMS  
AND SOFTWARE PROGRAM

**Richard Johansson**


`richard.johansson@gu.se`

## recap

- ▶ we saw how to define a function that computes a score for an input  $\mathbf{x}$  and an output  $\mathbf{y}$ :

$$\text{score}(\mathbf{x}, \mathbf{y}) = \sum_{i=1}^L \phi^e(\mathbf{x}, y_i) + \sum_{i=1}^L \phi^t(\mathbf{x}, y_{i-1}, y_i)$$

- ▶ the scoring function is “factorized” into **emission** scores  $\phi^e$  and **transition** scores  $\phi^t$

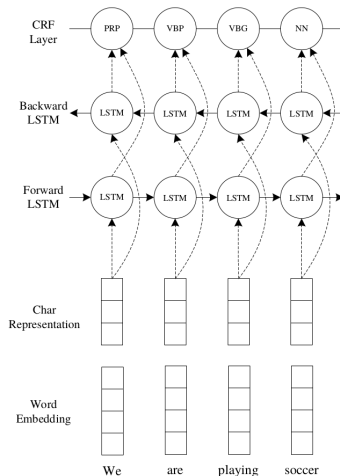


# training the scoring function

- ▶ there are many training algorithms that can train the scoring function  $\text{score}(\mathbf{x}, \mathbf{y})$
- ▶ some of them are generalizations of classification models:
  - ▶ perceptron  $\rightarrow$  structured perceptron ([Collins, 2002](#))
  - ▶ SVM  $\rightarrow$  structured SVM ([Tsochantaridis et al., 2005](#))
  - ▶ logistic regression  $\rightarrow$  **conditional random field**

# CRFs in NLP systems

- ▶ the CRF model was proposed by [Lafferty et al. \(2001\)](#) and has been popular in many NLP tasks since then
- ▶ in neural models, a CRF is typically used as the output layer ([Ma and Hovy, 2016](#))



# CRF and related models

probability model?		output type?	
		category	sequence
	<b>generative</b> $P(\mathbf{x}, \mathbf{y})$	naive Bayes, GMM	hidden Markov model
	<b>discriminative</b> $P(\mathbf{y} \mathbf{x})$	logistic regression	<b>conditional random field</b>

## CRF: basic definition (Lafferty et al., 2001)

- ▶ the CRF defines a probability of an output sequence  $\mathbf{y}$ , given an input sequence  $\mathbf{x}$ , as follows

$$P(\mathbf{y}|\mathbf{x}) = \frac{\exp \text{ score}(\mathbf{x}, \mathbf{y})}{\sum_{\mathbf{y}'} \exp \text{ score}(\mathbf{x}, \mathbf{y}')} = \frac{\exp \text{ score}(\mathbf{x}, \mathbf{y})}{Z(\mathbf{x})}$$

- ▶ to train the model, we minimize the negative log likelihood:

$$\mathcal{L}(\mathbf{X}, \mathbf{Y}) = - \sum_{i=1}^N \log P(\mathbf{y}_i | \mathbf{x}_i)$$

# CRF: basic definition (Lafferty et al., 2001)

- ▶ the CRF defines a probability of an output sequence  $\mathbf{y}$ , given an input sequence  $\mathbf{x}$ , as follows

$$P(\mathbf{y}|\mathbf{x}) = \frac{\exp \text{ score}(\mathbf{x}, \mathbf{y})}{\sum_{\mathbf{y}'} \exp \text{ score}(\mathbf{x}, \mathbf{y}')} = \frac{\exp \text{ score}(\mathbf{x}, \mathbf{y})}{Z(\mathbf{x})}$$

- ▶ to train the model, we minimize the negative log likelihood:

$$\mathcal{L}(\mathbf{X}, \mathbf{Y}) = - \sum_{i=1}^N \log P(y_i | x_i)$$

- ▶ but the probability is a giant softmax!
  - ▶ the sum in the denominator goes over all possible sequences!
  - ▶ the giant sum is called the **partition function**  $Z(\mathbf{x})$

# requirements for training and prediction

- ▶ for **predicting** or **decoding**, we need to compute the top-scoring output sequence  $\hat{\mathbf{y}}$  for a given input  $\mathbf{x}$ :

$$\hat{\mathbf{y}} = \arg \max_{\mathbf{y}} P(\mathbf{y}|\mathbf{x})$$

- ▶ for **training**, we need to compute the log likelihood, including the log of the partition function:

$$\log P(\mathbf{y}|\mathbf{x}) = \text{score}(\mathbf{x}, \mathbf{y}) - \log Z(\mathbf{x})$$



## finding the highest-probability sequence

- ▶ the partition function  $Z(\mathbf{x})$  does not depend on  $\mathbf{y}$
- ▶ so we don't need  $Z(\mathbf{x})$  to find the top-scoring output:

$$\begin{aligned}\arg \max_{\mathbf{y}} P(\mathbf{y}|\mathbf{x}) &= \arg \max_{\mathbf{y}} \frac{\exp \text{score}(\mathbf{x}, \mathbf{y})}{Z(\mathbf{x})} \\ &= \arg \max_{\mathbf{y}} \exp \text{score}(\mathbf{x}, \mathbf{y}) \\ &= \arg \max_{\mathbf{y}} \text{score}(\mathbf{x}, \mathbf{y})\end{aligned}$$


- ▶ we've seen that we can use the **Viterbi** algorithm to solve this

# the forward algorithm

- ▶ the **forward** algorithm computes

$$\log Z(\mathbf{x}) = \log \sum_{\mathbf{y}} \exp \text{ score}(\mathbf{x}, \mathbf{y})$$

- ▶ it is a special case of the **sum-product algorithm** for graphical models
- ▶ it uses a dynamic programming approach similar to Viterbi
- ▶ it computes log sums instead of maximizing




## side note about working in the log domain

- ▶ when computing  $\log Z(\mathbf{x})$  in the forward algorithm, we need to work in the **log domain** to avoid numerical overflows
- ▶ this is common when implementing probabilistic models in general
  - ▶ multiplying probabilities  $\Rightarrow$  summing log probabilities
  - ▶ summing probabilities  $\Rightarrow \log \sum \exp$  with log probabilities
- ▶ in PyTorch, `torch.logsumexp` computes  $\log \sum \exp$  in a numerically stable way

## the forward algorithm: dynamic programming

- ▶ assume we have computed the log-sum-exp for all paths in the previous step
  - ▶ let  $\alpha_{i-1,j}$  be the log-sum-exp of the scores of all paths ending in label  $j$  at position  $i-1$
- ▶ then we compute the  $\alpha$  scores in step  $i$  as follows:


$$\alpha_{ij} = \phi^e(y_i) + \log \sum_k \exp [\alpha_{i-1,k} + \phi^t(y_{i-1}, y_i)]$$



## the forward algorithm: dynamic programming

- ▶ assume we have computed the log-sum-exp for all paths in the previous step
  - ▶ let  $\alpha_{i-1,j}$  be the log-sum-exp of the scores of all paths ending in label  $j$  at position  $i-1$
- ▶ then we compute the  $\alpha$  scores in step  $i$  as follows:


$$\alpha_{ij} = \phi^e(y_i) + \log \sum_k \exp [\alpha_{i-1,k} + \phi^t(y_{i-1}, y_i)]$$



## the forward algorithm: dynamic programming

- ▶ assume we have computed the log-sum-exp for all paths in the previous step
  - ▶ let  $\alpha_{i-1,j}$  be the log-sum-exp of the scores of all paths ending in label  $j$  at position  $i-1$
- ▶ then we compute the  $\alpha$  scores in step  $i$  as follows:

$$\alpha_{ij} = \phi^e(y_i) + \log \sum_k \exp [\alpha_{i-1,k} + \phi^t(y_{i-1}, y_i)]$$




- ▶  $\log Z(\mathbf{x})$  is the  $\alpha$  score for the dummy end token

# CRF implementations for PyTorch

- ▶ PyTorch does not include a CRF module out of the box
- ▶ a couple of implementations:
  - ▶ **pytorch-crf**: a standalone CRF module
  - ▶ **AllenNLP**: the CRF module is part of a larger library
  - ▶ **NCRF++**: also part of a library


## sequence labeling: summary of approaches

- ▶ why have we spent so much time on sequence models?
- ▶ a fairly simple task, but we could explore different approaches to solving **structured prediction problems**





## summary: greedy sequential models



### ► pros:


- a bit more flexible, no need to design special algorithms
- since we reduce the task to classification, we can rely on standard machinery for classification

### ► cons:

- algorithm is greedy and does not necessarily find the highest-scoring sequence
- because prediction is sequential, we can't use any information about future predictions

## summary: global scoring approaches

$$\text{score}(\mathbf{x}, \mathbf{y}) = \sum_{p \in \text{parts}} \text{part-score}(\mathbf{x}, p) = \sum_{i=1}^L \phi^e(\mathbf{x}, y_i) + \sum_{i=1}^L \phi^t(\mathbf{x}, y_{i-1}, y_i)$$



- ▶ **pros:**
  - ▶ a bit more accurate (generally)
  - ▶ finds the globally best solution, does not go to a dead end
- ▶ **cons:**
  - ▶ need specialized algorithms for training and prediction
  - ▶ computational complexity may be higher

## exercise 2

- ▶ we will continue our NER experiments

Manchester	United	will	return	to	the	United	States
<b>B-ORG</b>	<b>I-ORG</b>	<b>O</b>	<b>O</b>	<b>O</b>	<b>O</b>	<b>B-LOC</b>	<b>I-LOC</b>

- ▶ we will investigate autoregressive models and CRFs

- ▶ Goldberg doesn't have a chapter on sequence tagging
  - ▶ chapter 19 includes a bit, but also some parts that are more relevant when we discuss parsing
- ▶ Eisenstein's chapter 7 covers CRF and related models
  - ▶ you can skim or skip the parts on HMMs and structured perceptrons and SVMs
- ▶ [Reimers and Gurevych \(2017b\)](#) gives a good overview of different engineering choices
  - ▶ ...and more in another paper ([Reimers and Gurevych, 2017a](#))

# state of the art for sequence labeling tasks

- ▶ Ruder collects state-of-the-art results for many NLP tasks:  
<https://nlpprogress.com/>
- ▶ examples of some of the tasks we have discussed:
  - ▶ part-of-speech tagging
  - ▶ named entity recognition
  - ▶ semantic role labeling
- ▶ quite English-centric. . .

## references

- M. Collins. 2002. Discriminative training methods for hidden Markov models: Theory and experiments with perceptron algorithms. In *EMNLP*.
- J. Lafferty, A. McCallum, and F. Pereira. 2001. Conditional random fields: Probabilistic models for segmenting and labeling sequence data. In *ICML*.
- X. Ma and E. Hovy. 2016. End-to-end sequence labeling via bi-directional LSTM-CNNs-CRF. In *ACL*.
- N. Reimers and I. Gurevych. 2017a. Optimal hyperparameters for deep LSTM-networks for sequence labeling tasks. arXiv:1707.06799.
- N. Reimers and I. Gurevych. 2017b. Reporting score distributions makes a difference: Performance study of LSTM-networks for sequence tagging. In *EMNLP*.
- I. Tsochantaridis, T. Joachims, T. Hofmann, and Y. Altun. 2005. Large margin methods for structured and interdependent output variables. *JMLR* 6(Sep):1453–1484.